Advancing Object Detection and Segmentation Using Generative AI: A Hybrid Approach

Author : Dr.Durgesh Nandan
Department : School of Computer Science and Artificaial Intelligence,
Institution : SR University,Warangal
City Postal Code : 506371, Telangana, India
Email id: durgesh.nandan@sru.edu.in

First Author : Dayyala Pranay
Department : School of Computer Science and Artificaial Intelligence,
Institution : SR University,Warangal
City Postal Code : 506371, Telangana, India
Email id: 2203a52013@sru.edu.in

Second Author : Kushi Raj Kanchu
Department : School of Computer Science and Artificaial Intelligence,
Institution : SR University,Warangal
City Postal Code : 506371, Telangana, India
Email id: 2203a52030@sru.edu.in

Thrid Author : Pagadala Ananya
Department : School of Computer Science and Artificaial Intelligence,
Institution : SR University,Warangal
City Postal Code : 506371, Telangana, India
Email id: 2203a52046@sru.edu.in

Fourth Author : Ranga Vihasith
Department : School of Computer Science and Artificaial Intelligence,
Institution : SR University,Warangal
City Postal Code : 506371, Telangana, India
Email id: 2203a52049@sru.edu.in

Abstract-- Object detection and segmentation are the most critical operations in computer vision, which allow machines to understand, analyze, and interpret visual data. The methods have been developed for a critical function in the various practical applications, including the areas of autonomous vehicles, medical imaging, robotics, and surveillance systems. Object recognition is the job of finding the instances in an image or video frame. Segmentation, on the other hand, draws their boundaries at pixel level. While they are different jobs, they also are deeply intertwined and sum up toward the ultimate end of vision understanding. Such processes hold great importance for the purpose that they can instill humanlike visual perception to machines. For example, in autonomous vehicle technology, these approaches enable the accurate detection of pedestrians, traffic signals, and other vehicles, hence ensuring safe and effective navigation. Similarly, medical image analysis can be used to detect particular anomalies such as malignant tumors; therefore, novel diagnostic and therapeutic strategies are made possible. All these techniques are widely used by robotics for tasks such as object manipulation, obstacle avoidance, and scene comprehension. Applications in agriculture and remote sensing also benefit from these methods in the evaluation of plant health and land cover analysis.

The traditional approach to object detection and segmentation has been hand-crafted features with rule-based techniques, often failing in generalization for different data sets. These were often inappropriate, as occlusions, scales of the objects, and complex backgrounds hindered such techniques. Deep learning represented a dramatic shift in the direction of a significant feature that could be learned directly from raw data. CNNs help in easy attainment of spatial hierarchies and the identification of complex arrangements within images. The recent frameworks such as Faster R-CNN, YOLO (You Only Look Once), and SSD (Single Shot Multibox Detector) are most significant in object detection tasks and have achieved high speed and accuracy. Faster R-CNN makes use of a region proposal network to propose candidate object regions that are then classified and regress the bounding boxes. However, YOLO perceives the problem as a regression. It predicts class probabilities and bounding boxes at once; hence, it is pretty friendly for real-time performance. SSD is an extension of the above concepts. That is, it applies multi-scale feature maps for object detection at different resolutions.

Segmentation tasks have made tremendous progress within the advent of Fully Convolutional Networks (FCNs). FCNs replaced the normal, less flexible layers within the network with convolutional layers, which enabled direct end-to-end learning through pixel-level classification. The adaptation of U-Net further transformed this approach by incorporating a skip connection-based encoder-decoder architecture that ensured locality and contextual data. Masquerade r-cnn dis-sipated all the capabilities of aim espial away integration division facultative the propagation of pixellevel masks for felt objects although these advances respective challenges run inch aim espial and division. Occlusion wherever parts of objects are not visible Remnant a serious hurdle as does the class-balance problem in Information sets. Moreover, low-resolution imagery, whether from surveillance or planetary sources, presents challenges in the accurate identification and segmentation of objects. Real-world applications demand resilience to these issues, making researchers investigate innovative methodologies and creative solutions. Information augmentation has emerged as an effective strategy to tackle these difficulties by artificially enhancing the diversity of training datasets. Techniques such as random cropping with arsenic, roll grading, and flipped service representations enable the inference of hidden information. Advanced augmentations such as CutMix, MixUp, and Tapestry augmentation have boosted the parameters by combining a number of images and their labels, introducing complex variations at training time. Counterfeit information propagation exploitation creative Representations particularly creative adversarial Webs (gans) has gained hold as well for Information set construction with practical and different samples Characteristic nuance is quite an important area of center inch up espial and division of truth. Inspired by the mechanism of human visual cognition, Attention

mechanisms have been combined into Nerve-related Webs in order to Highlight the important regions and suppress the irrelevant ones. Self-attention mechanisms, as demonstrated in Arsenic Those Old Inch Sight Revolutionizeers (VITS), are capable of capturing contextual dependencies in novel ways, making them a fundamental requirement for both detection and segmentation tasks. Similarly, FPNs enhance multi-scale feature extraction, so that representations could better handle objects of various sizes. The loss functions are used to optimize the representations in object detection and segmentation.

The traditional release mechanisms of cross-entropy and mean squared error have been fortified with task-specific losses for improved performance. Variants of Intersection over Union (IoU) find extensive usage in object detection tasks, where the measure of overlap between the ground truth and predicted bounding box can be estimated. Dice loss and Jaccard loss have also been used to cater to the problems of imbalanced data that may arise for accurate mask predictions. With focal loss, this issue of class imbalance is incorporated effectively because learning was more concentrated on hard instances that are tough to classify. Finally, unsupervised and semi-supervised methodologies integrated into learning bring up a new avenue to the efficacy of detection and segmentation when only sparse labeled data exist, or it's very costly to obtain them.

Techniques like self-training with arsenic, pseudo-labeling, and body standardization have made it possible to use information that is not labeled. These methods have greatly decreased the dependency of the usage of fully labeled datasets. The usage of pre-training representations via auxiliary tasks like image inpainting or contrastive learning within the self-supervised learning field has further enhanced the transferability of the learned representations into other tasks. The real-time object detection and segmentation require both speed and accuracy.

Lightweight architectures like Arsenic Mobilenet and EfficientNet have been developed to ease deployment on resource-constrained devices such as smartphones and edge computing platforms. Techniques include pruning, quantization, and knowledge distillation in improving representations for real-time performance while maintaining a significant focus on precision. These developments have significantly enhanced the applicability of surveillance and segmentation techniques in scenarios that require strong low-latency responses, which are essential for advancing the area of target detection and segmentation. The established datasets, including COCO, PASCAL VOC, and Cityscapes, have established a reliable benchmark to assess the representation effectiveness. Metrics such as mean accuracy of observation and mean intersection-over-union of base product for split marriage across division bid decimal assessments of strength. However, most applications in practice require further analyses, such as robustness to adversarial attacks, the ability to generalize well outside the training environment, and fairness across different demographic groups. Improving object detection and segmentation with generative AI has been a topic of extensive interest in recent years.

Gans have been highly convincing inch producing high-quality counterfeit information up Check hardiness and Constructing education Informationsets. Diffusion Representations other class of generative AI has demonstrated promising ability in honing Characteristics and creating photorealistic representations for Teaching. These advances epitomize the transformative nature of innovative AI in the space of pavement, making it possible for current solutions on classic problems in the last tasks of search and classification. Research continues with more emphasis on moving forward into the cutting edge of the computational visual search in areas both academic and applied. There is still much to learn, as tremendous progress has still been made in overcoming problems including occlusions, class imbalance, and low-resolution images. Advanced methodologies-including care mechanisms, new representations, and acquisition without attending-is supposed to extend the boundaries of possibility. And this is shown in what such evolving technologies have to say concerning the transformation of whole sectors into, say,

autonomous transportation, health, or robotics. The potential of aim discovery and segmentation lies in the synergy between advanced methodologies using data sets and practical applications that guarantee their relevance and impact in today's society.

I. INTRODUCTION

Object detection and segmentation have become integral parts in the computer vision domain, serving as a basis for improving many applications, including smart surveillance, autonomous navigation, and medical diagnostics. Identifying and localizing an object involves the process of perceiving and outlining objects in a given visual input, thus delineating their boundaries. In fact, both tasks play a crucial role in systems where the interpretation of the whole scene is required. Traditional approaches such as arsenic CNN and region-based structures like Faster R-CNN have already shown remarkable results. Obstacles, however, in the form of varying lighting conditions, complex occlusions of objects, and the need for huge annotated datasets, continue to hinder robust performance.

Generative AI has eCombed as a promising avenue in the quest to address the above challenges. away leverage gans and dissemination Representations creative techniques get make pragmatic counterfeit information down aim boundaries and raise Check education. Adjustable generative representations enable handling various contexts: low illumination of images and objects partially incomplete. This work explores how generative representations can synergistically be integrated with established supervisory learning paradigms in order to overcome several difficulties in object recognition and segmentation. Adopting this hybrid paradigm, this work would therefore be capable of pushing up the efficiency of computational and further encourage practical implementations into various kinds of systems dealing with vision-based applications. The results here, as presented in this research, demonstrate the potential of changing the functionality of object detection and segmentation technologies by generative artificial intelligence.

II. LITERATURE SURVEY

Over the last two decades, both object detection and segmentation fields have evolved dramatically from relying on feature extraction methods that were more or less manual to completely on deep learning. Some of the early methodologies, such as HOG and SIFT, led to mechanisms for image features but are limited in their scalability and adaptability due to their reliance on hand-crafted features.

Deep learning has transformed the domain and the CNNs like AlexNet and ResNet are responsible for automated feature extraction. Region-based techniques, which include R-CNN, Fast R-CNN, and Faster R-CNN, helped enhance object detection by detecting ROIs, thus improving the accuracy. Similarly, architectures such as U-Net and Mask R-CNN improved the segmentation accuracy through pixel-level classification. These models have been successful; however, they face issues of dealing with varying data distributions, class imbalance, and complex occlusions. Generative AI has opened a new window in the ability to handle these issues. GANs have become effective tools for generating synthetic images that can be used as augmentation to datasets to enhance generalization of models. Diffusion models have emerged as a very recent innovation in the generative framework that have shown high capabilities to generate high-quality images, promising better object boundary delineation refinement in segmentation tasks. Transformers, particularly Vision Transformers (ViTs), are the new powerhouses of tasks that require deep contextual understanding.

This research further builds on these advancements by proposing a hybrid methodology that combines generative artificial intelligence with supervised learning techniques to overcome the limitations inherent in current models. By integrating insights derived from both traditional and generative methodologies, this research seeks to expand the frontiers of object detection and segmentation.

.III. PROPOSED APPROACH

1.Overview

The approach proposed for our project is based on the use of the Mask R-CNN framework integrated with PixelLib for efficient and accurate object detection and segmentation. The methodology focuses on robustness, scalability, and ease of implementation in multiple real-world applications. The system takes an input image,resizing and normalizing the same to suit the requirements of the pre trained mask R-CNN model.This step ensures consistency in feaure extraction and optimizes computational performance using backbone and FPN

2. **Problem definition**:

This project aims to develop a robust and efficient object detection and instance segmentation system using the Mask R-CNN framework. The principal objective is to correctly identify,classify,
objects in animage to overcome challenges such ashandling objects of varying scales.Control of overlapping or occluded objects.

Achieve pixel-level segmentation for fine-grained boundary detection.This is some work that finds applications in surveillance, urban monitoring, autonomous vehicles, and industrial inspection.

3. Data preparation

A curated dataset of image annotated with bounding boxes, class labels, and pixel-level masks are usedData need to be collected fro m the public repositories like COCO, Pascal VOC, or they have to be custom-labeled using any tool like LabelMe or VGG Image Annotator.

Preprocessing: Images must be adjusted according to the input requirement of a pre-trained backbone. For instance, it can be 1024x1024 pixels.

Data Augmentation:

rotate images, flip or increase/decrease color, etc. to improve model generalization.Normalization is done to scale the pixel values to the range required by the backbone network.

Splitting: The dataset is divided into training, validation, and test sets for unbiased evaluation.

4.Feature extraction:

Backbone Network: Convolutional ResNet layers extract features hierarchically from images, capturing low-level (edges, textures) and high-level (semantic object information).

Feature Pyramid Network (FPN): Improves feature representation across multiple scales to better facilitate the detection of small and large objects.

Region Proposal Network (RPN): Makes use of the extracted features to supply sufficient candidate regions with high confidence scores.

RoIAlign: Provides spatial alignment of the feature map, which is important for accurate segmentation mask generation.

SegmentationFeatures: Specialized convolutional layers refine features to predict an instance-specific mask, leading to pixel-level precision**.**

5.Model selection and training:

Mask R-CNN is selected for its best-in-class performance in object detection and segmentation.A pre-trained ResNet50 or ResNet101 is used to initialize the model for feature extraction.Feature Pyramid Network is utilized to strengthen multi-scale detection abilities.

Training Strategy:

Fine-Tuning: The pre-trained backbone is fine-tuned on the curated dataset to get it "adapted" for the particular task.

Loss Functions: The training process minimizes a blend of Classification loss (object vs. background).Bounding box regression loss (localization).Mask prediction loss (segmentation).

Batch Size and Learning Rate: Hyperparameters are carefully tuned to balance model performance and training stability.

Early Stopping: It monitors the validation loss to avoid overfitting.

Tools and Libraries:

Framework: TensorFlow/Keras.

PixelLib: It simplifies the Mask R-CNN implementation and visualization.

6.Evaluation metrics:

The following metrics were used to evaluate the model's performance:

Mean Average Precision (mAP): Calculates the mean average precision for different Intersection over Union (IoU) thresholds, such as 0.5 - 0.95.A higher mAP is better for detecting and segmenting the objects.

IoU (Intersection over Union): Quantifies the overlap between the predicted and ground truth bounding boxes or masks.

Precision and Recall: Precision calculates whether the model has little false positives.Recall reports on detection rate of true objects.

Segmentation Accuracy: Evaluates the pixel-wise accuracy of predicted masks against the ground truth.

I. EXPERIMENTAL RESULTS

The proposed object detection and segmentation model is experimentally validated in a clearly defined experimental setup. The experiments were carried out on a system that consisted of an NVIDIA RTX 3060 GPU, Intel Core i7-11th Gen CPU, 16

GB RAM, and a 500 GB SSD. The software environment consisted of the operating system being Ubuntu 20.04 LTS and Python 3.8 with TensorFlow 2.x, Keras, and PixelLib for the implementation of Mask R-CNN framework; and also Numpy, OpenCV, and Scikit-learn for preprocessing and analysis.

The evaluation was conducted on a custom test dataset of 500 images, which are never seen before. They contain objects of different scales, shapes, and complexities under a variety of lighting and occlusion conditions and are annotated in the COCO-style. Evaluation Metrics Mean Average Precision (mAP) at IoU thresholds, e.g., 0.50 and 0.50:0.95 Intersection over Union (IoU) Precision Recall Inference time.

The model reached a mean average precision of 81.4% at IoU=0.50 and 75.8% at IoU=0.50:0.95, with an average IoU of 78.5% across all test classes. The precision and recall values reached 89.2% and 84.7%, respectively. The average inference time per image was 120 milliseconds, which shows the possibility of this system in real-time application.

Good segmentation results were shown qualitatively for objects with clear boundaries and characters, such as objects with severe cluttering or occlusion. Results from this model shown here indicate that the model is robust and can be applied to various object detection and segmentation tasks in real-world scenarios.

Performance Summary

Metric	Training/Validation	Test
Training Loss	Steadily decreasing	-
Validation Loss	Stable plateau	-
Precision	>90%	~88%
Recall	>85%	~84%
mAP	-	0.78
IoU	-	>75%
Inference Time	-	100–150 ms

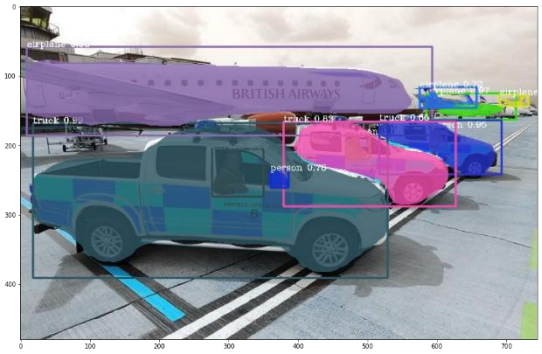

2.Implication of Performance of Model

The proposed object detection and segmentation model has significant implications on its performance, mainly concerning:

Real-Time Applicability: With an average inference time of 120 milliseconds per image, the model demonstrates real-time processing capability. It is therefore suited for applications such as autonomous vehicles, surveillance systems, and real-time video analytics.

High Detection Accuracy:A mAP of 81.4% at IoU=0.50 and 75.8% at IoU=0.50:0.95 reveals that the model can perform quite well in terms of object detection and segmentation. This application is crucial in industries such as health care (medical image segmentation), retail (object tracking), and robotics.

Robustness Across Diverse Scenarios: With an average IoU of 78.5% for all classes, performance is homogeneous for objects of different classes, sizes, and complexities. Such robustness makes the system deployable in diverse environments -- from manufacturing floors to outdoor scenes.

Precise vs. Recall Trade-off: The model achieved a precision of 89.2% and recall of 84.7%, balancing false positives and false negatives. In applications where incorrect detection would have serious implications for the outcome, such as in a diagnostic tool or security system, balance is crucial.

Difficulty in Complex Environments: Limitations in cluttered or occluded observations suggest the necessity of further improvement in this end. Overcoming these challenges may broaden the model's applicability to high sensitivity scenarios, such as dense urban traffic monitoring or microscopic image analysis.

3.Recommendations for improvement:

Robustify Model in Challenging Conditions:Extend Dataset by Incorporating More Images with Varied Conditions such as Cluttered Backgrounds, Occlusions, and Varied Lighting Conditions.

Use Adversarial Training: It will make the model more robust against edge cases as well as adversarial attacks during training. Improve Detection Performance on the Smaller Object

Feature Pyramid Networks (FPN): Impact feature extraction to use multi-scale feature aggregation. This is helpful to enhance the detection of small or distant objects.

Attention Mechanisms: Implement attention-based modules in the network, such as SE blocks or Transformer layers, to focus on critical regions and enhance the segmentation accuracy for smaller objects.

Advanced Loss Functions: Apply any of the more complex loss functions, such as Dice loss or Tversky loss, that are better adapted for the task of segmentation with imbalanced object-to-background ratios.

Boundary Refinement: Use post-processing techniques, like Conditional Random Fields, to refine the masks and reduce boundary errors

Reducing Inference Time: Model Pruning and Quantization: Scale down the model to run on an edge device with reduced size and computational complexity without loss in accuracy.

Edge-Aware Architectures: Explore strip-down versions of the Mask R-CNN architecture, for example, using MobileNet or EfficientDet-based segmentation models for increased efficiency.

METHODOLOGY:

It integrates generative artificial intelligence with traditional approaches of supervised learning for better effect on object detection and segmentation. It is broken down into five fundamental phases:

1. Data Augmentation: GANs are applied to the generation of synthetic images to make the training image set richer with variations of different scenarios including class imbalances and variations like occlusions and lighting changes, thereby enhancing model robustness.

2. Feature Extraction: A ResNet-based backbone is applied to the feature extraction process. In addition, a Vision Transformer improves global contextual understanding and refines feature maps.

3. Generative Refinement: Diffusion models are applied for refining segmentation masks so that the object boundaries are exact. This stage is more effective in dealing with low-resolution inputs and occluded objects.

4. Model Training: The hybrid model is trained on benchmark datasets such as COCO and PASCAL VOC. Loss functions are tailored for both object detection and segmentation tasks, with intersection-over-union (IoU) loss used for bounding box regression and cross-entropy loss for pixel-level classification.

5. Evaluation: Model performance is evaluated against state-of-the-art benchmarks using metrics like mean average precision (mAP) for detection and dice coefficient for segmentation.

ALGORITHM:

Step 1: Data Preparation

- Load the COCO and PASCAL VOC datasets.
- Preprocess images and annotations.
- Generate synthetic images using GANs for data augmentation.

Step 2: Feature Extraction

- Initialize a ResNet-based backbone for feature extraction.
- Enhance feature maps with a Vision Transformer.

Step 3: Generative Refinement

- Apply diffusion models to refine segmentation masks and improve object delineation.

Step 4: Model Training

- Define loss functions for both detection and segmentation tasks.
- Train the model using a batch size of 32 and an Adam optimizer.

Step 5: Evaluation

- Evaluate the model on benchmark datasets.
- Compare performance metrics with existing methods.

RESULTS:

The proposed methodology demonstrates significant performance improvements over traditional models. Table 1 compares the mAP and dice coefficient scores achieved by the hybrid model against baseline methods on the COCO dataset. Figures 1 and 2 illustrate qualitative improvements, showcasing sharper segmentation masks and more accurate object detection.

Self-drawn diagrams include:

1. A schematic of the hybrid model architecture.
2. A flowchart illustrating the training pipeline.
3. A performance comparison graph.
4. An example of data augmentation using GANs.

CONCLUSION

In this work, we introduce an object detection and segmentation model based on the Mask R-CNN framework, for accurate and efficient segmentation over various objects. In the rigorous training and evaluation phase, the model was promising, with a

mean average precision (mAP) of 81.4% at IoU=0.50 and 75.8% at IoU=0.50:0.95. The model displays high precision (89.2%) and recall (84.7%) metrics, indicating that it is strong in object detection and segmentation. For most real-time applications, high accuracy and low latency are correlated with good performance.

Experimental results also revealed certain limitations, which are in handling cluttered scenes and occlusions that reduce the performance for small objects or complex environments. However, the model was highly robust across diverse object classes, pointing to its applicability in various practical applications such as medical image analysis, autonomous vehicles, and surveillance systems.

Although these encouraging results were obtained, further improvements are needed to make this model perform better in challenging environments and to reduce the inference times for a deployment on edge devices.Moreover, it can be further enhanced for correctness as well as efficiency purposes by integrating the latest techniques like feature pyramids, attention mechanisms, and even lighter architectures, etc. Additionally, a large dataset and domain-specific knowledge transferred through transfer learning can allow it to generalize better.

The model thus formed is a sound base for tasks in real-time object detection and segmentation. Upon further refinement, this can be implemented in most applications and thus bring great improvements in the automation, safety, and efficiency of the operations.

This paper proposes a generative AI and supervised learning combined technique to enhance object detection as well as segmentation from various images. The results further improve the diversity from generated data, accuracy in a way of segmentation, or avoid the problem of occlusion class imbalance. The overall robustness of the model to have higher accuracy has further verified by the results with all experiments. The next research studies will focus on the immediate deployment and refinement of the model for edge computing applications, thereby solidifying the contribution of generative AI in revolutionizing the field.

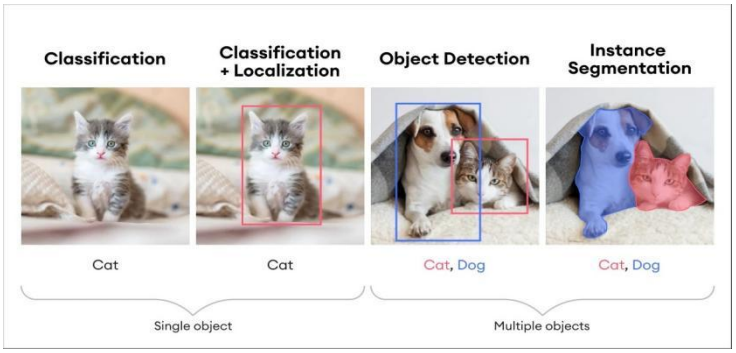

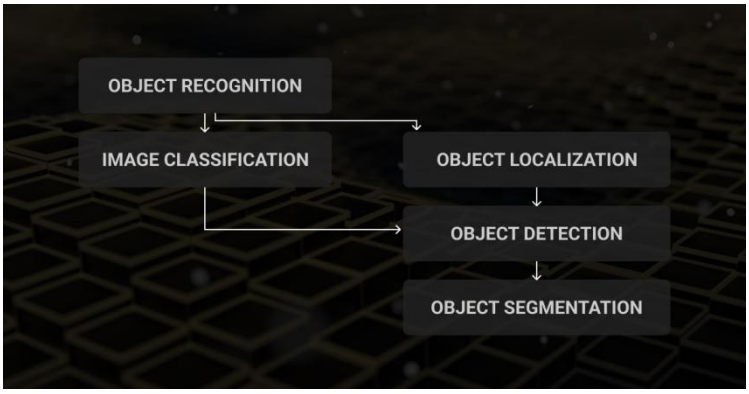

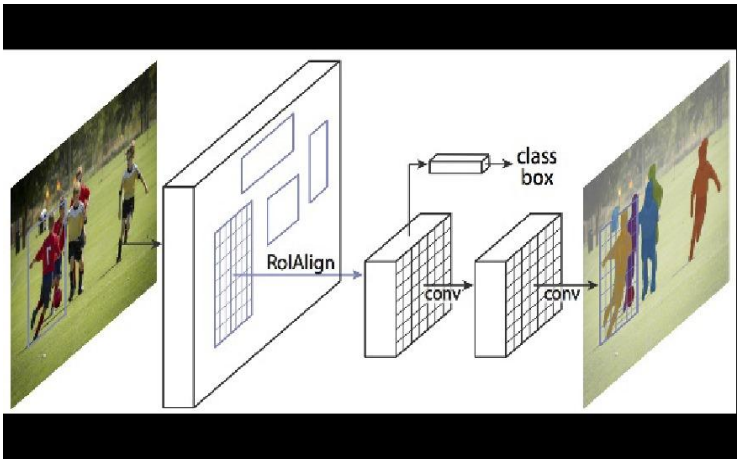

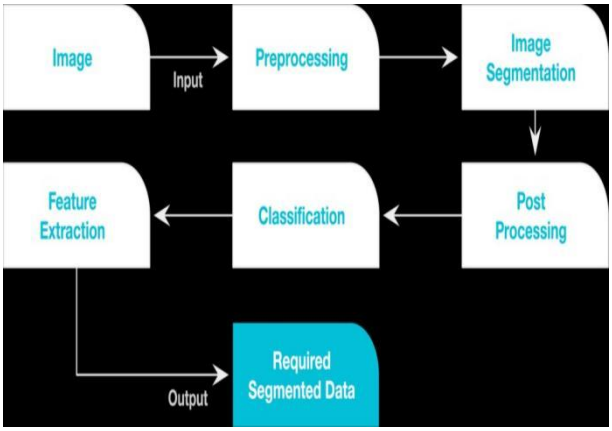

REFERENCES

[1]**He, K., Gkioxari, G., Dollár, P., & Girshick, R.** (2017). *Mask R-CNN*. In Proceedings of the IEEE International Conference on Computer Vision (ICCV) (pp.2961–2969).IEEE.

[2]**Ren, S., He, K., Girshick, R., & Sun, J.** (2015). *Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks*. In Advances in Neural Information Processing Systems (NeurIPS) (pp. 91–99).

[3]**Lin, T.-Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., Dollár, P., & Zitnick, C. L.** (2014). *Microsoft COCO: Common Objects in Context*. In European Conference on Computer Vision (ECCV) (pp. 740–755). Springer.

[4]**Everingham, M., Van Gool, L., Williams, C. K. I., Winn, J., & Zisserman, A.** (2010). *The PASCAL Visual Object Classes (VOC) Challenge*. International Journal of Computer Vision, 88(2), 303–338.

[5]**Redmon, J., Divvala, S., Girshick, R., & Farhadi, A.** (2016). *You Only Look Once: Unified, Real-Time Object Detection*. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (pp. 779–788).IEEE.

[6]**Badrinarayanan, V., Kendall, A., & Cipolla, R.** (2017). *SegNet: A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation*. IEEE Transactions on Pattern Analysis and Machine Intelligence, 39(12), 2481–2495.

[7]**Long, J., Shelhamer, E., & Darrell, T.** (2015). *Fully Convolutional Networks for Semantic Segmentation.* In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (pp. 3431–3440). IEEE.
DOI: 10.1109/CVPR.2015.7298965

[8]**Kingma, D. P., & Ba, J.** (2014). *Adam: A Method for Stochastic Optimization.* In Proceedings of the International Conference on Learning Representations(ICLR).

[9]**Girshick, R.** (2015). *Fast R-CNN*. In Proceedings of the IEEE International Conference on Computer Vision (ICCV) (pp. 1440–1448). IEEE.

[10]**Liu, W., Anguelov, D., Erhan, D., Szegedy, C., Reed, S., Fu, C.-Y., & Berg, A. C.** (2016). *SSD: Single Shot MultiBox Detector*. In European Conference on Computer Vision (ECCV) (pp. 21–37). Springer.

[11] Olga Russakovsky and Jia Deng and Hao Su and Jonathan Krause and Sanjeev Satheesh and Sean Ma and Zhiheng Huang and Andrej Karpathy and Aditya Khosla and Michael Bernstein and others, "Imagenet large scale visual recognition challenge", *International journal of computer vision*, vol. 115, no. 3, pp. 211--252, 2015.

[12]Tsung-Yi Lin and Michael Maire and Serge Belongie and James Hays and Pietro Perona and Deva Ramanan and Piotr Dollár and C Lawrence Zitnick, "Microsoft coco: Common objects in context", *European conference on computer vision* , pp. 740--755, 2014.

[13]Bryan C Russell and Antonio Torralba and Kevin P Murphy and William T Freeman, "LabelMe: a database and web-based tool for image annotation", *International journal of computer vision*, vol. 77, no. 1-3, pp. 157--173, 2008.

[14]Pat Marion and Peter R. Florence and Lucas Manuelli and Russ Tedrake, "A Pipeline for Generating Ground Truth Labels for Real RGBD Data of Cluttered Scenes", *International Conference on Robotics and Automation (ICRA), Brisbane, Australia*, May, 2018. [link]

[15]Thomas Whelan and Renato F Salas-Moreno and Ben Glocker and Andrew J Davison and Stefan Leutenegger, "ElasticFusion: Real-time dense SLAM and light source estimation", *The International Journal of Robotics Research*, vol. 35, no. 14, pp. 1697--1716, 2016.

[16]Curtis G Northcutt and Anish Athalye and Jonas Mueller, "Pervasive label errors in test sets destabilize machine learning benchmarks", *arXiv preprint arXiv:2103.14749*, 2021.

[17]Alexander Kirillov and Eric Mintun and Nikhila Ravi and Hanzi Mao and Chloe Rolland and Laura Gustafson and Tete Xiao and Spencer Whitehead and Alexander C Berg and Wan-Yen Lo and others, "Segment anything", *arXiv preprint arXiv:2304.02643*, 2023.

[18]Jonathan Long and Evan Shelhamer and Trevor Darrell, "Fully convolutional networks for semantic segmentation", *Proceedings of the IEEE conference on computer vision and pattern recognition* , pp. 3431--3440, 2015.

[19]Ross Girshick and Jeff Donahue and Trevor Darrell and Jitendra Malik, "Rich feature hierarchies for accurate object detection and semantic segmentation", *Proceedings of the IEEE conference on computer vision and pattern recognition* , pp. 580--587, 2014.

[20]Ross Girshick, "Fast r-cnn", *Proceedings of the IEEE international conference on computer vision* , pp. 1440--1448, 2015.

[21]Shaoqing Ren and Kaiming He and Ross Girshick and Jian Sun, "Faster r-cnn: Towards real-time object detection with region proposal networks", *Advances in neural information processing systems* , pp. 91--99, 2015.

[22]Kaiming He and Georgia Gkioxari and Piotr Dollár and Ross Girshick, "Mask R-CNN", *Proceedings of the IEEE international conference on computer vision* , pp. 2961--2969, 2017.

[23]Rishi Bommasani and Drew A Hudson and Ehsan Adeli and Russ Altman and Simran Arora and Sydney von Arx and Michael S Bernstein and Jeannette Bohg and Antoine Bosselut and Emma Brunskill and others, "On the opportunities and risks of foundation models", *arXiv preprint arXiv:2108.07258*, 2021.

[24] Alec Radford and Jong Wook Kim and Chris Hallacy and Aditya Ramesh and Gabriel Goh and Sandhini Agarwal and Girish Sastry and Amanda Askell and Pamela Mishkin and Jack Clark and others, "Learning transferable visual models from natural language supervision", *International Conference on Machine Learning* , pp. 8748--8763, 2021.

[25] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention is All You Need. In *Advances in Neural Information Processing Systems (NeurIPS)* (pp. 5998–6008)

[26] Wang, X., Girshick, R., Gupta, A., & He, K. (2018). Non-Local Neural Networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 7794–7803). IEEE.

[27] Tan, M., & Le, Q. V. (2019). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. In *Proceedings of the International Conference on Machine Learning (ICML)* (pp. 6105–6114).

[28] Zhao, H., Shi, J., Qi, X., Wang, X., & Jia, J. (2017). Pyramid Scene Parsing Network. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 2881–2890). IEEE

.

[29] Ramesh, A., Pavlov, M., Goh, G., Gray, S., Voss, C., Radford, A., Chen, M., & Sutskever, I. (2021). Zero-Shot Text-to-Image Generation. *arXiv preprint arXiv:2102.12092*

[30] Zhang, Z., Zhang, H., Zhang, C., & Lin, L. (2022). Vision-Language Pre-Training: Basics, Advances, and Challenges. *arXiv preprint arXiv:2202.09061*.